

eXtreme Gradient Boost

-

what calculus enthusiasts love the most

WhyR??


Tianqi Chen

- Github: [github](#)
- Twitter: [Twitter](#)
- Professor in the Machine Learning Department and Computer Science Department at Carnegie Mellon University.
- Research
 - Apache TVM
 - **XGBoost**
 - Apache MXNet



What is it?

- One of the most liked machine learning algorithms in Kaggle.
- Teams with this algorithm win the competition.
- It can be used for supervised learning tasks such as regression and classification.



Gradient boosting



Gradient boosting for regression

```
In [7]: from sklearn.tree import DecisionTreeRegressor  
  
tree_reg1 = DecisionTreeRegressor(max_depth=2, random_state=42)  
tree_reg1.fit(X, y)
```

```
Out[7]: DecisionTreeRegressor(max_depth=2, random_state=42)
```

```
In [34]: y2 = y - tree_reg1.predict(X)  
tree_reg2 = DecisionTreeRegressor(max_depth=2, random_state=42)  
tree_reg2.fit(X, y2)
```

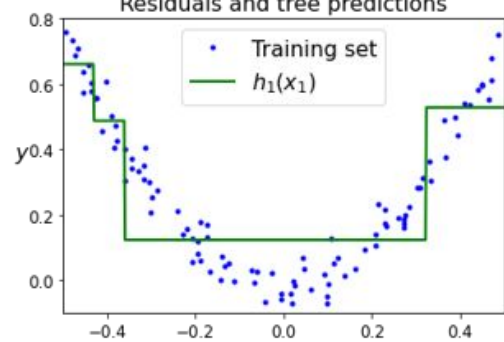
```
Out[34]: DecisionTreeRegressor(max_depth=2, random_state=42)
```

```
In [35]: y3 = y2 - tree_reg2.predict(X)  
tree_reg3 = DecisionTreeRegressor(max_depth=2, random_state=42)  
tree_reg3.fit(X, y3)
```





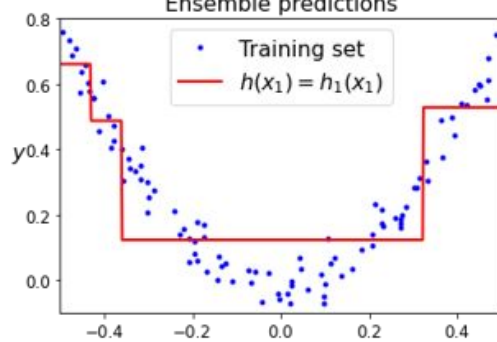
```
Out[35]: DecisionTreeRegressor(max_depth=2, random_state=42)
```

```
In [37]: y_pred = sum(tree.predict(X_new) for tree in (tree_reg1, tree_reg2, tree_reg3))
```

Residuals and tree predictions



Ensemble predictions



But what is a relationship between XGBoost and gradient?

Loss function $L(y, F(x)) = (y - F(x))^2/2$

We want to minimize $J = \sum_i L(y_i, F(x_i))$ by adjusting $F(x_1), F(x_2), \dots, F(x_n)$.

Notice that $F(x_1), F(x_2), \dots, F(x_n)$ are just some numbers. We can treat $F(x_i)$ as parameters and take derivatives

$$\frac{\partial J}{\partial F(x_i)} = \frac{\partial \sum_i L(y_i, F(x_i))}{\partial F(x_i)} = \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} = F(x_i) - y_i$$

So we can interpret residuals as negative gradients.

$$y_i - F(x_i) = -\frac{\partial J}{\partial F(x_i)}$$

$$F(x_i) := F(x_i) + h(x_i)$$

$$F(x_i) := F(x_i) + y_i - F(x_i)$$

$$F(x_i) := F(x_i) - 1 \frac{\partial J}{\partial F(x_i)}$$

$$\theta_i := \theta_i - \rho \frac{\partial J}{\partial \theta_i}$$

For regression with **square loss**,

residual \Leftrightarrow negative gradient

fit h to residual \Leftrightarrow fit h to negative gradient

update F based on residual \Leftrightarrow update F based on negative gradient

XGBoost Parameters

- General parameters:
 - **booster** - possible values:
 - gblinear
 - gbtree
 - gbdart (“Dropouts meet Multiple Additive Regression Trees”)
- Booster parameters
- Learning task parameters:
 - **objective**
- Command line parameters:
 - only used in the CLI version of XGBoost

Booster parameters

- eta (learning_rate)
- gamma (min_split_loss)
- min_child_weight
- max_depth
- alpha (L2) & lambda (L1)

More info about available parameters can be found in the [XGBoost documentation](#)

Preparing data for XGBoost

- Before applying XGBoost, we have to **convert** all **data** into **numeric** type
 - Label Encoding (e.g. Species Category in Iris dataset)

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
1	1	5.1	3.5	1.4	0.2	Iris-setosa
80	80	5.7	2.6	3.5	1.0	Iris-versicolor
150	150	5.9	3.0	5.1	1.8	Iris-virginica



	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_encoded
1	1	5.1	3.5	1.4	0.2	0
80	80	5.7	2.6	3.5	1.0	1
150	150	5.9	3.0	5.1	1.8	2

Preparing data for XGBoost

- Before applying XGBoost, we have to convert all data into numeric type
 - Label Encoding - code

```
12 df <- data.frame(read_csv('Iris.csv'))
13 categories <- df[['Species']]
14 lbl <- LabelEncoder$new()
15 lbl$fit(categories)
16 encoded <- lbl$transform(categories)
17 df_encoded <- data.frame(df[, 1:length(df)-1], encoded)
```

Preparing data for XGBoost

- Before applying XGBoost, we have to convert all data into numeric type
 - Label Encoding
 - One Hot Encoding (e.g. in dataset regarding breast cancer cases)

	x.40.49.	x.premeno.	x.15.19.	x.0.2.	x.yes.	x.3.	x.right.	x.left_up.	x.no.	x.recurrence.events.
1	'50-59'	'ge40'	'15-19'	'0-2'	'no'	'1'	'right'	'central'	'no'	'no-recurrence-events'
2	'50-59'	'ge40'	'35-39'	'0-2'	'no'	'2'	'left'	'left_low'	'no'	'recurrence-events'
3	'40-49'	'premeno'	'35-39'	'0-2'	'yes'	'3'	'right'	'left_low'	'yes'	'no-recurrence-events'

	x.3..2.	x.3..3.	x.right..left.	x.right..right.	x.left_up..central.	x.left_up..left_low.
1	0	0	0	1	1	0
2	1	0	1	0	0	1
3	0	1	0	1	0	1
	x.left_up..left_up.	x.left_up..right_low.	x.left_up..right_up.	x.left_up.nan	x.no..no.	
1	0	0	0	0	1	
2	0	0	0	0	1	
3	0	0	0	0	0	
	x.no..yes.	x.recurrence.events..no.	x.recurrence.events..recurrence.events.			
1	0	1	0			
2	0	0	1			
3	1	1	0			

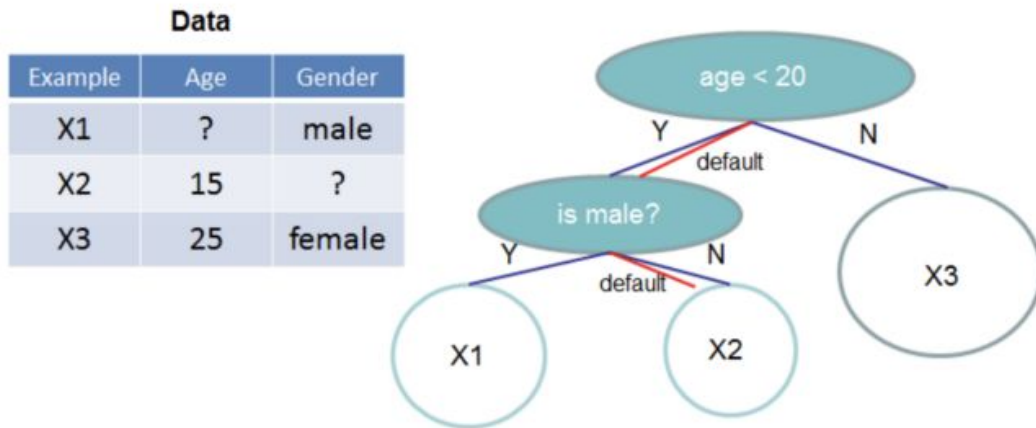
Preparing data for XGBoost

- Before applying XGBoost, we have to convert all data into numeric type
 - Label Encoding
 - One Hot Encoding - code

```
19 # One Hot Encoding Breast Cancer
20 library(caret)
21
22 df2 <- data.frame(read_csv('breast-cancer.csv'))
23
24 dummy <- dummyVars("~ .", data=df2)
25 final_df2 <- data.frame(predict(dummy, newdata=df2))
```


Preparing data for XGBoost

- Before applying XGBoost, we have to convert all data into numeric type
- Missing values **are dealt with automatically** by XGBoost during model training



Preparing data for XGBoost

- Before applying XGBoost, we have to convert all data into numeric type
- Missing values are dealt with automatically by XGBoost during model training
- XGBoost works with data in ***matrix*** type - ***DataFrame*** type is not supported
 - In order to convert Data Frame to a matrix, you can use this R function:

```
# Converting Data Frame to a Numeric Matrix  
data.matrix(df)
```

xgboost package

since 2016

90K download per month

1. Can solve regression, classification and ranking problems,
2. Two solvers are included: linear model and tree based model.
3. Can automatically do parallel computation on Windows and Linux.

Package usage

INSTALLATION

```
install.packages('xgboost')
```

library(xgboost)

```
# zbiór z pakietu xgboost  
data(agaricus.train, package='xgboost')  
data(agaricus.test, package='xgboost')  
train <- agaricus.train  
test <- agaricus.test
```

CLASSIFICATION

```
bstSparse <- xgboost(data = train$data, label = train$label,  
max.depth = 2, eta = 1, nthread = 2, nrounds = 2,  
objective = "binary:logistic")  
pred <- predict(bstSparse, test$data)  
as.numeric(head(pred > 0.5))
```

REGRESSION

```
bstSparse <- xgboost(data = train$data, label = train$label,  
max.depth = 2, eta = 1, nthread = 2, nrounds = 2,  
objective = "reg:squarederror")  
pred <- predict(bstSparse, test$data)
```

Related literature

1. [Short documentation with basic usage.](#)
2. [Detailed documentation.](#)
3. [Chen, Tianqi, and Carlos Guestrin. "Xgboost: A scalable tree boosting system." 2016.](#)
4. [Machine learning with XGBoost tutorial.](#)

Bibliography

<https://www.doit-intl.com/xgboost-or-tensorflow/>

<https://towardsdatascience.com/xgboost-is-not-black-magic-56ca013144b4>

<https://xgboost.readthedocs.io/en/stable/faq.html>

<https://xgboost.readthedocs.io/en/stable/parameter.html>

http://www.chengli.io/tutorials/gradient_boosting.pdf