WB

Kacper Grzymkowski
Faculty of Mathematics and Information Science
Warsaw University of Technology
Poland
`kacper.grzymkowski.stud@pw.edu.pl`

Dominik Kędzierski
Faculty of Mathematics and Information Science
Warsaw University of Technology
Poland
`dominik.kedzierski.stud@pw.edu.pl`

Jakub Piwko
Faculty of Mathematics and Information Science
Warsaw University of Technology
Poland
`jakub.piwko.stud@pw.edu.pl`

May 12, 2022

ABSTRACT

Machine Learning has become powerful tool for data scientists in wide variety of application. Particularly interesting problem which shows efficiency of artificial intelligence algorithms is real estate appraisal. This paper examines and compares popular or recently developed tree-based models including Random Forest, XGBoost, LightGBM and Catboost for this regression task. Research is conducted on data describing numerous aspects of residential homes in Ames, Iowa, United States. We perform exploratory data analysis to investigate set characteristics. Preprocessing pperations were selected in order to achive more intuitive way of variables coding. Next step is building models and tuning hyperparameters. Different methods for finding the most optimal parameters set were used for each model. We concentrate on models diagnosis, which determine the best models regarding minimizing of RMSE and MAE scores but also their susatinability. Explainable artificial intelligence methods were used to unveil decision-making process of best models.

Keywords Machine Learning · Explainable Artificial Intelligence · Real estate · Property appraisal

1 Introduction

...

2 Data

2.1 Data Description

Data analysed in our article was originally collected by Dean De Cock for use in data science education. It is publicly available through an advanced regression techniques machine learning competition on kaggle.com. The data describes 2919 dwellings located in Ames, Ihoa, United States. It is split between a training set of 1460 observations and a test set of 1459 observations. There are 79 feature columns describing various aspects of the dwelling such as its lot area, quality of amenities or the distance to the nearest railroad. The label is the sale price of the real estate.

2.2 Conclusions from Exploratory Data Analysis

Early analysis found that there are significant counts of missing values within the dataset. However, upon further inspection, those missing values were tied closely to the structure of the data. Certain features, such as *Basement Quality* or *Swimming Pool Quality* do not make sense when those features are missing. As such, we decided to impute the missing data by creating new categories representing a missing feature.

What is interesting in our data set is presence of many attributes that hold some kind of rating. For example, *Overall Quality* provides information about the overall quality of dwelling. We could not find a concrete methodology with which this data was collected. It is possible for these ratings to be subjective. Additionally some of these variables were in a categorical format, and some already encoded into an integer format. We categorized all feature columns into 47 discrete, 16 ordered discrete, and 16 continuous features.


Figure 1: Relationship between dependant and target variable

Next step was analysing some one-dimensional and multi-dimensional dependencies by visualising them in histograms, box plots and scatter plots. As a result, we found some variables that seem to impact appraisal of real estate as shown on plots 2.2. For example, built year, type of foundation and area showed some relation with price of dwellings, so we hope that these columns will be significant in predicting prices.
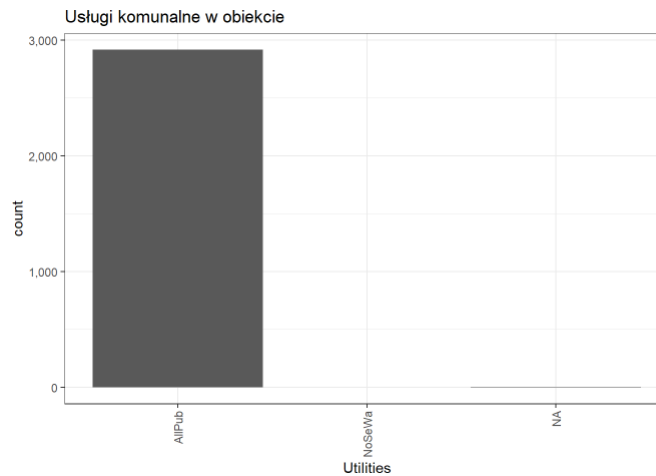

Figure 2: Type of utilities count

Exploratory Analysis outcome has also implied that in some columns there is a huge dominance of one value, while other types of observations are negligible as presented on plot 2.2. There are also variables related to extra features like pool or tennis courts and only a small part of dwellings presented in frame own them. This is a hint to encode this kind of columns in a different way or just not pay that much attention to them in a process.


Figure 3: Correlation heatmap

Our final stage was to look more broadly at continuous variables, so we created correlation heatmap 2.2. We haven't found any particularly high correlations between variables. But we observed that negative correlations are being outnumbered, which can be related to specification of our data.

3 Model selection

Data are explored and properly processed. The next crucial step of our research is building models. Studying this particular problem of property appraisal is a great opportunity to test some of lately popular tree-based models. In our research we are focusing on following algorithms:

- Random Forest
- XGBoost
- LightGBM
- Catboost

All of them rely on ensemble of many decision trees to improve forecasting, but use different types of variables coding, individual methods for finding split points and gradient boosting in case of 3 last types. All of these factors raise expectations about their performance and begs the question, which one is going to be the best in case of real estate price prediction.

Models training was preceded by process of writing code that were to wrap building procedures for each model type into single function. As said before, our models differs if it comes to preparation of data or model construction, so our research group divided and worked on automation of training models for each algorithm separately.

After reaching the goal, parameters tuning is to be conducted. Every model type has its own and in some cases exceptionally long list of parameters to choose from. Decision was made to perform and adjust tuning for every model kind and then diagnose them using proven metrics and visualisations, with emphasis of root mean square error.

3.1 CatBoost

CatBoost hyperparameters were selected by a combined grid and random search approach. The search was performed over parameters exported from the `catboost` Prokhorenkova et al. [2017] package into the `caret` Kuhn [2008] API, namely: `iterations`, `depth`, `l2_leaf_reg`, `learning_rate`, `rsm`. Initial grid search found `learning_rate` and `iterations` to be the most important hyperparameters. These findings were used to refine the search space and a random search was performed afterwards. Results of tuning are presented in table 1.

RMSE Train	RMSE Test	depth	iterations	learning_rate	l2_leaf_reg	rsm
3312.41	24471.06	5	800	0.1	0.05	0.75
3906.97	24483.43	5	700	0.1	0.05	0.75
5604.79	24563.16	5	500	0.1	0.05	0.75
3280.82	25975.99	4	1400	0.1	0.01	0.90
397.50	24430.67	6	1200	0.1	0.01	0.90

Table 1: Catboost hyperparameter tuning results.

3.2 LightGBM

Attempts of LightGBM training started with the moment of preprocessing. Examination of different operations preceding learning has shown that the best scores are reached by selecting only automatic handling data with parameters `autofactor` and `forceconvert` set to TRUE. These variables are available in function implemented beforehand. Decision was made that the best way to find the most sustainable set of training parameters is to create our own grid of hyperparameters and pass them randomly to models. Tuning principles for lightGBM were taken into consideration during selection of values. As a result, optimal random grid of parameters was was constructed, including: `max_depth`, `num_leaves`, `bagging_fraction`, `feature_fraction`, `learning_rate`, `num_iterations`, `min_data_in_leaf`, `lambd_l1`. Then again, after fitting 200 models, we examined the RMSE and MAE scores on train and test sets and chose best 5 regarding these metrics. Table 2 shows final values for each parameter of our top 5 models.

	max_depth	num_leaves	num_iterations	learning_rate	bagging_fraction	feature_fraction	lamba_l1	min_data_in_leaf
1	3	2	242	0.06619538	0.572464	0.4015881	1.37272030	31
2	4	2	182	0.18479657	0.5591560	0.7143656	0.02852181	39
3	4	2	54	0.14623834	0.8445833	0.6481415	1.24955739	27
4	3	2	347	0.17079440	0.5978079	0.4352913	0.72251183	33
5	3	2	88	0.14964913	0.9925450	0.6002572	0.40267346	21

Table 2: Parameters of best LightGBM models

In table 3 RMSE of models for training and test set in corresponding order was collected. Surprisingly, models generated with lightGBM are very stable. The difference between scores on train and test set is diminutive. It turned out that this kind of tree model is the best if to comes to sustainability in our research. Unfortunately, metrics alone are not so impressive. The threshold of 40000 on test data was barely crossed.

	RMSE_train	RMSE_test
1	39869.42	39968.16
2	39101.91	40065.58
3	40663.25	41350.35
4	37729.43	41060.23
5	39295.34	41700.98

Table 3: RMSE for best LightGBM models

3.3 XGBoost

For training XGBoost models, firstly we performed preprocessing on our data – the same we did while exploring data. It included imputation of missing values and coding of variables concerning quality or conditions rate.

For parameters tuning caret package mentioned before was used. By default, list of hyperpyrameters to perform matching included: `eta`, `max_depth`, `gamma`, `colsample_bytree`, `min_child_weight`, `subsample`, `nrounds`.

Over 200 models were trained using cross-validation and 'gbtree' option. In next step 40 best were picked. Evaluation was carried out with RMSE and MAE metrics taken into account. We were trying to choose models that not only performed well on test set, but also had low difference between metrics on test and train sets to cope with overfitting. Then chosen parameters were fitted to functions implemented beforehand and again models performance was diagnosed

by measuring RMSE and MAE. Finally, selection of 5 best in our opinion was completed. Table 4 shows corresponding parameters values.

	eta	max_depth	gamma	min_child_weight	subsample	colsample_bytree	nrounds
1	0.04462844	3	3.580246	1	0.5339151	0.6035825	305
2	0.05525418	3	1.062662	10	0.9131714	0.6054747	305
3	0.19592307	2	8.526370	16	0.5217781	0.6626087	81
4	0.20411763	2	3.240579	1	0.6497563	0.6905933	405
5	0.08603091	2	3.549911	0	0.6369859	0.4167207	163

Table 4: Parameters of best XGBoost models

Error scores for best XGBoost models are contained in table 5. They reached pretty low values on a test set and significantly lower when predicted train values, which implies poor stability. Sadly, none of other models have reduced the gap between scores more satisfactorily. Although, XGBoost performance is acceptable and these top 5 models can compete with others.

	RMSE_train	RMSE_test
1	12005.92	28219.85
2	13583.38	30127.77
3	22388.00	34033.15
4	10474.84	29589.40
5	24614.30	36718.85

Table 5: RMSE for best XGBoost models

3.4 Random Forest

To build the random forest model as in previous models we started with data preprocessing. We removed columns containing missing data. We binarized data frame columns that were characters and had only two unique values inside each other. In the last step we performed one hot encoding on given dataframe. OHE was performed on character columns with more than two unique categories only.

To have some reference in further tuning the hyperparameters we started by building our model on the default parameters. We obtained RMSE values on it of 14025.22 for the training set and 64501.26 for the test set.

Hyperparameter tuning was performed using the random search method. For this purpose, we prepared a grid of the following hyperparameters:

- `num.trees` - number of trees in model, 400 to 1000 increasing in 20 steps
- `max.depth` - Maximal tree depth, 5 to 100 increasing in 5 steps
- `min.node.size` - Minimal node size, 1 to 10 increasing in 1 steps
- `splitrule` - Splitting rule, extratrees or vaiance

We trained 100 models from which we selected the top 5 models based on the average RMSE values on the training and test sets. They obtained the following RMSE values:

	RMSE_train	RMSE_test
1	49542.65	38349.32
2	43751.19	43696.35
3	43422.84	44024.12
4	30349.42	55754.17
5	30553.19	55836.22

3.5 Comparing models

Once we had selected the top 5 models for random forest, xgboost, lightgbm and catboost we decided to compare them to each other. To do this we created a point chart of the RMSE values for each of the 20 selected models.

Figure 4: Comparing RMSE

The analysis of the models turned out to be fairly easy. It is straightforward to see that random forest turned out to be the worst model. For this model both the training and test RMSE values turned out to be very large compared to the other models. The second worst model is LightGBM, the only advantage of this model is that the training and test RMSE values are very close to each other, but they are still very large. Third turned out to be XGBoost. The best model on the other hand is CatBoost. Models built with CatBoost have a very small training RMSE, test RMSE is much larger, but it is still much smaller than for random forest, xgboost and lightgbm.

3.6 Selection and diagnosis of the best model

Of all the models, we chose the best one with a training RMSE of 397.50 and a test RMSE of 24430.67. This is the CatBoost model with the following parameters:

- depth - 6

- n_estimators - 1200

- learning_rate - 0,1

- l2_leaf_reg - 0,01

- colsample_bylevel - 0,9

We started the diagnostics of the selected model by calculating the MAE (Mean absolute error) value. It is 254.4542 for the training set and 14987.81 for the test set.

Figure 5: Percentage difference - train

As you can see, the difference between the true and predicted values for the training set is almost always less than 1


Figure 6: Percentage difference - test

For the test set, the difference is significantly larger, but except for a few outlier observations, for the vast majority the difference is only a few percent. We can also see from the graph that our model more often tends to overestimate rather than underestimate property values.

References

Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. Catboost: unbiased boosting with categorical features, 2017. URL `https://arxiv.org/abs/1706.09516`.

Max Kuhn. Building predictive models in r using the caret package. *Journal of Statistical Software, Articles*, 28(5): 1–26, 2008. ISSN 1548-7660. doi:10.18637/jss.v028.i05. URL `https://www.jstatsoft.org/v028/i05`.

Cankun Wei, Meichen Fu, Li Wang, Hanbing Yang, Feng Tang, and Yuqing Xiong. The research development of hedonic price model-based real estate appraisal in the era of big data. *Land*, 11(3):334, February 2022. doi:10.3390/land11030334. URL `https://doi.org/10.3390/land11030334`.

Geoffrey K. Turnbull and Arno J. van der Vlist. After the boom: Transitory and legacy effects of foreclosures. *The Journal of Real Estate Finance and Economics*, February 2022. doi:10.1007/s11146-021-09882-w. URL `https://doi.org/10.1007/s11146-021-09882-w`.

Alex van de Minne, Marc Francke, and David Geltner. Forecasting US commercial property price indexes using dynamic factor models. *Journal of Real Estate Research*, 44(1):29–55, December 2021. doi:10.1080/08965803.2020.1840802. URL `https://doi.org/10.1080/08965803.2020.1840802`.

Dieudonné Tchuente and Serge Nyawa. Real estate price estimation in french cities using geocoding and machine learning. *Annals of Operations Research*, February 2021. doi:10.1007/s10479-021-03932-5. URL `https://doi.org/10.1007/s10479-021-03932-5`.

Víctor Hugo Masías, Mauricio Valle, Fernando Crespo, Ricardo Crespo, Augusto Vargas Schüler, and Sigifredo Laengle. Property valuation using machine learning algorithms: A study in a metropolitan-area of chile. 01 2016.

Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. Lightgbm: A highly efficient gradient boosting decision tree. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, NIPS'17, page 3149–3157, Red Hook, NY, USA, 2017. Curran Associates Inc. ISBN 9781510860964.